



# CAP Theorem Notes

## 1. Introduction to CAP Theorem

- Proposed by Eric Brewer (2000).
  - In the presence of a network partition, a distributed system can guarantee at most two of the following three:
    1. **Consistency (C)**
    2. **Availability (A)**
    3. **Partition Tolerance (P)**
  - Real distributed systems must choose trade-offs depending on requirements.
- 

## 2. Properties Explained

### 2.1 Consistency (C)

- Every read sees the most recent write, like a single up-to-date database.
- Equivalent to: "all users get the same answer at the same time."
- Example: Banking system must always show the correct balance.

### 2.2 Availability (A)

- The system always responds to requests (success or failure) even if some nodes are down.
- It doesn't guarantee the freshest data, only that you get some response.
- Example: Shopping website shows cached prices when DB is lagging.
- Doubt: "If the system is down, how can it respond?"
  - Availability means the system as a whole keeps serving because at least some replicas are up (not the crashed one).

### 2.3 Partition Tolerance (P)

- System continues to operate despite network failures (messages between nodes may be lost).

- You can't avoid partitions in large networks → must tolerate them.
  - Doubt: "Partition tolerance sounds like availability — both say system keeps working if one node is down?"
    - Difference:
      - Availability → node/server crashes but others still serve.
      - Partition tolerance → network split between groups of nodes, both groups may still serve requests independently.
- 

### 3. Trade-Offs in CAP

#### 3.1 CA (Consistency + Availability)

- Works only when no partitions exist.
- If a partition happens, one property must be dropped.
- Rare in real-world distributed systems.

#### 3.2 CP (Consistency + Partition Tolerance)

- Guarantees:
  - Always consistent
  - Survives partitions
- Sacrifices: Availability (some requests rejected).
- Examples: ZooKeeper, HDFS NameNode
- Behavior:
  - If a node (like C) is partitioned away, it refuses writes (✗ try later).
  - A+B continue serving and stay consistent.
  - When partition heals, C syncs from A+B.
- Doubt: "If someone writes to A or B while C is cut off, won't that break consistency?"
  - No, because C rejects writes. Only A+B accept. After healing, C catches up → consistency preserved.

#### 3.3 AP (Availability + Partition Tolerance)

- Guarantees:
    - Always available
    - Survives partitions
  - Sacrifices: Consistency (may diverge → eventual consistency).
  - Examples: DynamoDB, Cassandra
  - Behavior:
    - If C is partitioned away, C' (replica) may take over.
    - Both sides accept writes (A+B vs C').
    - After healing, conflicts resolved via:
      - Last-write-wins (timestamps)
      - Vector clocks (causality tracking)
      - CRDTs (auto-merge data structures)
  - Doubt: "Why can't A+B sync with C' during partition?"
    - Because partition means no communication path exists. Only after healing can they exchange logs and resolve.
  - Doubt: "After healing, how do A+B decide whether to sync with C or C'?"
    - Cluster consensus protocol (Raft/Paxos, Gossip, etc.) decides:
      - If A+B had majority → C/C' syncs from A+B.
      - If C' was promoted → A+B sync with C'.
      - Conflict resolution may merge histories.
- 

## 4. Timeline Example

- Partition Event:
  - T1 → A+B and C' lose contact.
  - T2 → A+B continue writes. C' may also accept writes.
- Healing Event:
  - T3 → Network restored.
  - T4 → Sync happens:

- CP system → C' discards its writes, catches up from A+B.
- AP system → Conflicts reconciled (timestamps/vector clocks).

## 5. Real-World Examples

- CP Systems: HDFS, Zookeeper, Google Spanner (strong consistency).
- AP Systems: DynamoDB, Cassandra, CouchDB (eventual consistency).
- CA Systems (rare, mostly theoretical at large scale): Relational DBs in a single-node cluster.

### ✓ Final Takeaway

- You cannot avoid partitions in real networks.
- So every real distributed system is either CP or AP.
- Choice depends on business needs:
  - Banking → CP (better to reject a transaction than show wrong balance).
  - E-commerce → AP (better to show slightly stale inventory than reject orders).

## Quick Comparison Table

Property	CA	CP	AP
Consistency	✓ (only without partitions)	✓	✗ (eventual)
Availability	✓ (only without partitions)	✗ (may reject)	✓
Partition Tolerance	✗	✓	✓
Typical Use Cases	Single-node RDBMS	Metadata stores, coordination (ZooKeeper, Spanner)	High-throughput, user-facing (Cassandra, DynamoDB)

## CP vs AP Decision Checklist

- ☐ Strong invariants required at all times? Choose **CP**

- ☐ Okay to return slightly stale data if system stays up? Choose **AP**
  - ☐ Can clients retry on write failures? **CP** is viable
  - ☐ Need write availability during partitions? Prefer **AP**
  - ☐ Human-facing latency critical over perfect freshness? **AP**
  - ☐ Cross-region, globally distributed with strict ordering needs? Likely **CP** with consensus
  - ☐ Conflict-free data types or easy merge semantics available? **AP** gets easier
  - ☐ Regulatory or financial correctness constraints? Bias to **CP**
- 

## Flashcards: CAP Theorem

- ▼ What does CAP stand for?
  - Consistency, Availability, Partition Tolerance
- ▼ Who proposed the CAP theorem and when?
  - Eric Brewer, 2000
- ▼ What is the core statement of CAP?
  - In the presence of a network partition, a system can provide at most two of Consistency, Availability, and Partition Tolerance.
- ▼ Define Consistency in CAP.
  - Every read sees the most recent write, as if from a single up-to-date copy.
- ▼ Define Availability in CAP.
  - Every request receives a non-error response, without guarantee that it contains the most recent write.
- ▼ Define Partition Tolerance in CAP.
  - The system continues operating despite network faults that prevent some nodes from communicating.
- ▼ How do Availability and Partition Tolerance differ?
  - Availability concerns serving requests despite node failures. Partition tolerance concerns continuing operation despite network splits

between groups of nodes.

- ▼ In a partition, what must a real distributed system choose between?
  - Consistency and Availability
- ▼ What does a CP system sacrifice during partitions?
  - Availability, by rejecting or delaying some requests.
- ▼ What does an AP system sacrifice during partitions?
  - Immediate consistency, allowing temporary divergence.
- ▼ Give two examples of CP systems.
  - ZooKeeper, Google Spanner
- ▼ Give two examples of AP systems.
  - Cassandra, DynamoDB
- ▼ What is eventual consistency?
  - Replicas may diverge temporarily but converge to the same state after no new updates and successful communication.
- ▼ Name three conflict resolution strategies in AP systems.
  - Last-write-wins, vector clocks, CRDTs
- ▼ In CP, what happens to writes on a minority partition?
  - They are rejected; clients are asked to retry later.
- ▼ In AP, why can't partitions synchronize during the split?
  - There is no communication path; sync occurs only after healing.
- ▼ How does a cluster decide which side to follow after healing?
  - Consensus or membership protocols determine leadership or majority; others reconcile to it.
- ▼ When might you choose CP over AP?
  - When strong invariants and correctness are critical, such as financial systems.
- ▼ When might you choose AP over CP?
  - When high availability and low latency are more important than immediate consistency, such as user-facing catalogs.